

Relative Performance Feedback and Redistribution Preferences: Evidence from an Online Experiment*

Isabella Valencia

Paris School of Economics

valenciaisabellap@gmail.com

May 2026

Abstract

This paper examines how relative performance feedback affects redistribution preferences and beliefs about effort versus luck in generating inequality. In an online experiment, participants complete a code-recognition task and are randomly assigned to one of three conditions: no feedback, told they performed *above* a randomly matched peer, or told they performed *below*. They then act as impartial spectators and redistribute a fixed sum between two workers whose earnings reflect their performance. Participants told they outperformed their peer score 1.23 points higher on an effort-attribution scale (0–10) than those told they underperformed ($p < 0.05$) and redistribute \$0.52–0.55 less to the lower-earning worker ($p < 0.10$). The no-feedback group sits between the two treatment arms; neither its belief score nor its redistribution differs significantly from either treated group. A mediation analysis shows that roughly 50 percent of the redistribution effect is explained by the shift in effort–luck beliefs, consistent with self-serving attribution as the channel linking feedback to redistribution.

Keywords: redistribution, relative performance feedback, self-serving attribution, fairness, online experiment

JEL codes: C91, D31, D63, D91

*This experiment builds on the design of Almås et al. (2020) and is inspired by the research question of Cassar and Klein (2019).

1 Introduction

Free-market economies inherently generate unequal outcomes, and while modern democratic societies rely on redistribution to mitigate these disparities, there is persistent disagreement over which inequalities are considered legitimate (Alesina and Giuliano, 2011; Piketty, 2014). A central dimension of this debate concerns the *source* of inequality: outcomes attributed to effort are typically viewed as more deserving than those arising from luck. This intuition is formalised in the accountability framework, which models redistributive preferences as judgments about deservingness and predicts lower redistribution when inequality reflects effort rather than chance (Konow, 2000; Cappelen et al., 2007). The framework finds broad empirical support: beliefs about the determinants of income are consistently among the strongest predictors of redistribution preferences (Piketty, 1995; Alesina and Angeletos, 2005; Bénabou and Tirole, 2006).

Yet a key difficulty in this debate is that, in practice, the source of inequality is rarely clear-cut. Individual success reflects a mix of effort, innate ability, family background, and circumstance, and it is often hard to separate these contributions, even for the person directly involved (Piketty, 2014). When effort and luck cannot be cleanly distinguished, beliefs about which is responsible do not simply reflect the available evidence; they are also shaped by personal experience and motivated reasoning (Kunda, 1990). Since these beliefs drive redistribution preferences, whatever shapes them also shapes redistribution demand.

Deffains et al. (2016) argue that the experience of success and failure shapes redistribution preferences through the self-serving bias: individuals who succeed at a task attribute their outcome to effort and favour less redistribution, while those who fail attribute their outcome to circumstance and shift toward more. Since wealthier individuals have typically experienced more success, this bias systematically pushes them to view inequality as earned rather than lucky, creating a divide between rich and poor not just in material interests but also in fairness views. As Deffains et al. (2016) put it, wealthy individuals “do not only oppose redistribution because they expect to be net payers, but also because the self-serving bias systematically shifts their fairness principles.” Identifying this channel outside a controlled setting is difficult, however: in observational data one cannot tell whether successful individuals attribute outcomes to effort because of cognitive bias or because they genuinely worked harder. Isolating the self-serving bias therefore requires exogenous variation in comparative signals, holding actual performance constant while shifting only what individuals learn about their relative standing. This paper asks whether such a signal — being told one outperformed or underperformed a peer — is enough to shift beliefs about the role of effort in generating inequality and, through that channel, redistribution preferences.

This paper provides experimental evidence on this mechanism. Participants complete a performance task and are randomly assigned to one of three conditions: (i) *no feedback* about relative standing; (ii) told they performed *above* a randomly matched peer (*above* group); or (iii) told they performed *below* a matched peer (*below* group). Subsequently, all participants act as impartial spectators and decide how to distribute a fixed monetary sum between two workers whose inequality arose from differential performance.

The key findings are three. First, participants told they outperformed a peer redistribute less in the merit scenario: \$1.91 on average, versus \$2.44 in the below group and \$2.12 in the no-feedback group. The above–below gap of \$0.52–0.55 is marginally significant ($p < 0.10$) across both specifications. Second, the same group scores 1.23 points higher on an effort-attribution scale (0–10) than the below group ($p < 0.05$), consistent with self-serving belief updating; the joint test across all three groups also rejects equality of means ($p = 0.039$, ANOVA). **Third, neither treatment arm individually differs significantly from the no-feedback group; the significant contrast is the above–below gap. This is consistent with both signals pushing beliefs in opposite directions from a common baseline, rather than one side driving the effect asymmetrically. Together, these results suggest that receiving a comparative signal — in either direction — shifts effort attribution and redistribution relative to receiving no signal, with the two treated groups diverging significantly from each other.**

These results are consistent with the self-serving attribution account of redistribution developed by Deffains et al. (2016): individuals who perform well relative to others update their beliefs about the role of effort in generating inequality and become less redistributive. Our design extends their framework by using randomly assigned comparative feedback, providing clean causal identification and isolating the relative rather than absolute performance channel. One notable difference is that Deffains et al. (2016) find SSB effects on both sides — success reduces redistribution and failure increases it — whereas our evidence is concentrated on the winning side. The absence of a significant response among participants told they underperformed suggests that the SSB in our setting is more asymmetric: learning one beat a peer is enough to trigger meritocratic updating, but learning one lost does not produce a symmetric shift toward redistribution. The paper most directly relates to Almås et al. (2020), from whose spectator design we adapt the merit redistribution task, and to Cassar and Klein (2019), whose focus on how performance comparisons shape distributive preferences motivates our research question. Our contribution is to provide clean experimental identification of the direction and asymmetry of this effect using

randomly assigned relative feedback. [REVISAR: confirmar asimetría con test lineal below vs. no-feedback antes de finalizar.]

The rest of the paper is structured as follows. Section 2 reviews the related literature. Section 3 describes the experimental design. Section 4 presents the results. Section 5 concludes.

2 Related Literature

This paper contributes to three strands of literature.

Beliefs about inequality and redistribution. A key insight from Piketty (1995) is that beliefs about social mobility and the role of effort versus luck determine preferences for redistribution in equilibrium. Alesina and Angeletos (2005) provide empirical evidence: Americans attribute poverty more to individual failure than Europeans, which partly explains the cross-country gap in redistribution. Bénabou and Tirole (2006) model how self-serving motivated reasoning sustains two equilibria — a “belief in a just world” and a more luck-oriented view — with different steady-state redistribution levels. Our experiment directly tests the mechanism by which an informational shock (relative performance feedback) updates beliefs and behaviour.

Impartial spectator and fairness views. Cappelen et al. (2007) and Cappelen et al. (2013) use spectator designs to classify subjects into fairness types (libertarian, meritocratic, egalitarian) and show that a majority of individuals prefer some redistribution away from inequality that arises from luck rather than effort. We adapt their spectator design to a merit scenario and ask whether third-party redistribution decisions are affected by information about the spectator’s own relative performance — a test of whether the spectator role is truly impartial.

Self-serving attribution and social comparisons. Self-serving attribution refers to the tendency to attribute successes internally (effort) and failures externally (luck or circumstance) (Miller and Ross, 1975; Mezulis et al., 2004). In an economic context, Babcock and Loewenstein (1995) document that parties in bargaining inflate self-assessments of a fair outcome in a self-serving direction. Closer to our design, Deffains et al. (2016) show that individuals who experience unfair treatment become more supportive of redistribution. Cassar and Klein (2019) show that framing one’s own performance as a failure (relative to a counterfactual) shapes distributive preferences; our design extends this insight by using

randomly assigned comparative feedback, allowing clean causal identification. We add to this literature by focusing on the *comparative* and *asymmetric* dimension: the same absolute performance level can produce different attribution and redistribution depending on whether one learns one stands above or below a peer, and the effect appears concentrated on the winning side.

3 Experimental Design

The experiment was conducted online using JATOS and hosted the free, ESCoP-sponsored server Mindprobe. The whole design builds on the spectator design of Almås et al. (2020). It involved two distinct roles: *workers* and *spectators*.

Workers ($n = 2$) were recruited through Prolific and completed the code-recognition task under real monetary incentives. They were paid a base fee of \$4 plus a performance-based bonus of up to \$8, for a maximum possible payout of \$12 each. They were informed in advance that their earnings would initially depend on their relative performance against each other—the higher-scoring worker would receive \$8 while the other would receive \$0—and that a third party (a spectator) would subsequently decide whether and how much to redistribute between them.

Spectators ($n = 134$) were recruited through personal networks (friends, family, and colleagues from two workplaces), completed the full experiment online in approximately 10 minutes, and received no monetary compensation. The experiment consisted of two parts: a code-recognition task and a redistribution decision task in which they acted as impartial spectators. Spectators were informed before making their redistribution decisions that the workers were real individuals who had completed the same task and would receive actual monetary payouts determined by one randomly selected spectator’s allocation. The impartial-spectator design thus removes material self-interest on the spectator side while keeping real stakes for the workers, ensuring that decisions have genuine consequences and giving the design greater external validity.

3.1 Part 1: Code-Recognition Task

Participants (both workers and spectators) completed a code-recognition task consisting of one practice round of 20 seconds followed by five scored rounds of 45 seconds each. At the start of each round, a three-digit target number (drawn from 100–999) was displayed at the top of the screen. Below it, a grid of three-digit numbers appeared; approximately 15 percent of the codes in the grid matched the target (around 48 instances per round), and

the rest were distractors. Participants had to identify and check every instance of the target as quickly as possible. The grid was generated randomly at the start of each round, so the target and distractor codes changed every 45 seconds. To ensure a consistent interface experience, the grid dimensions were adapted to screen size: desktop users saw a 20×16 grid (320 codes), while mobile users saw a 7×46 grid (322 codes), with the same number of target instances in both cases. Participants were explicitly told about the scoring system: they earned 1 point for each correct identification, lost 1 point for each false alarm, and received no penalty for a missed target. The maximum possible score across the five rounds was around 200 points; sample means range from 90 to 95 points across treatment groups (Table 1).¹

Just after finishing, participants answered three pre-feedback self-report questions: (i) effort exerted (slider, 0–10); (ii) guessed score (numeric, range -200 to 200); and (iii) relative performance prediction (“I would perform better than or the same as the other participant” vs. “I would perform worse”).

3.2 Treatments: Relative Performance Feedback

After the self-report, participants were randomly assigned to one of three feedback conditions.

1. **No feedback.** Participants were informed they could proceed to Part 2 without seeing their score or any comparison information.
2. **Above.** Participants were told: “You performed above your randomly matched peer.”
3. **Below.** Participants were told: “You performed below your randomly matched peer.”

The feedback signal was generated through a dynamic matching procedure. Before the main data collection, ten pilot participants completed the task, providing an initial distribution of scores and establishing a current best and worst performer. As each subsequent spectator completed the task, those randomly assigned to the *above* condition were matched with the current lowest scorer in the spectator pool; those assigned to the *below* condition were matched with the current highest scorer. Reference scores were updated continuously as the experiment progressed. Note that spectators were compared with other spectators in this pool and the matching was indeed random. This design guarantees genuine, non-deceptive feedback: every participant in the above condition truly outperformed their assigned reference peer, and every participant in the below condition truly underperformed

¹This marked the end of the task for the two workers. Before exiting, they were reminded that one spectator, chosen at random, would decide how to split their earnings, and that this decision would be paid out as their bonus. Spectators continued to Part 2, as described below.

theirs. Participants who could not receive their assigned treatment — above-condition participants who scored below the current worst, or below-condition participants who scored above the current best — were excluded from the analysis. At debriefing, participants were shown their actual score and informed of the matching procedure, so they could understand why they had received the message they did.

3.3 Part 2: Distribution Decisions

After the (non)feedback screen, all participants acted as impartial spectators and made a redistribution decision adapted from Almås et al. (2020).

Redistributive Decision in a Merit Scenario

Participants read that two real participants — the two workers who had previously completed the same task — were identified as Worker A and Worker B. Worker A had received a larger initial payment because they performed better; Worker B received nothing because they performed worse. Participants were asked to decide whether and how to redistribute the given endowment between the two workers. The decision was presented as a set of discrete options covering the full range from keeping full inequality to reversing it.

After the redistribution decision, participants completed a short demographics survey (age, gender, and education) followed by the key belief measure of this study. They were asked to respond to the following question using a slider:

“To what extent do you think the initial income (before any redistribution) of the participants in the previous task was determined by effort versus luck?”

The slider ranged from 0 (“Completely determined by luck”) to 10 (“Completely determined by effort”). This variable (`effort_luck`) is the hypothesised mediator between relative performance feedback and redistribution preferences.

3.4 Key Outcome and Mediator Variables

- **Redistribution (merit):** USD allocated to Worker B in the merit scenario. Primary outcome variable.
- **Effort–luck belief (`effort_luck`):** Score on the 0–10 slider described above. Hypothesised mediator.

3.5 Sample

Spectators were recruited through personal networks: friends, family, and employees at two workplaces. The analysis sample consists of $N = 145$ spectators after excluding non-engagers (zero score and zero clicks). The sample is predominantly Colombian (76.5%), with the remainder drawn from other countries, primarily in Europe. As this is a convenience sample, findings should be interpreted with caution regarding external validity and may not generalise to other populations in terms of magnitudes; treatment assignment was, however, random within this sample, so internal validity is preserved.

Table 1 presents descriptive statistics by treatment group. Groups are balanced on all pre-determined covariates: age ($p = 0.94$), gender ($p = 0.57$), task score ($p = 0.71$), nationality ($p = 0.32$), and education ($p = 0.52$), confirming that randomisation was successful. The nationality variable shows the above group with a slightly higher share of Colombians (83%) compared to the no-feedback (70%) and below (77%) groups, though this difference is comfortably non-significant ($p = 0.32$) and attributable to chance given the small sample size.

The effort–luck belief score differs significantly across groups ($p = 0.039$), which is itself a key outcome; importantly, this measure is elicited *after* the feedback treatment, so it cannot confound treatment assignment.

Table 1: Descriptive Statistics by Treatment Group


	Told: Above	No Feedback	Told: Below	p -value
Redistribution – merit (\$)	1.91 (1.35)	2.12 (1.39)	2.44 (1.38)	0.179
Effort–luck belief (0–10)	7.42 (2.36)	6.96 (2.03)	6.22 (2.55)	0.041
Task score	94.53 (27.81)	92.50 (26.37)	89.85 (27.89)	0.706
Age	34.77 (13.69)	35.84 (14.97)	35.21 (14.93)	0.935
Female (%)	61.70 (49.14)	62.00 (49.03)	70.83 (45.93)	0.567
Colombian (%)	82.98 (37.99)	70.00 (46.29)	77.08 (42.47)	0.319
Education (%)				0.519
Less than bachelor	23.40 (42.80)	18.00 (38.81)	10.42 (30.87)	
Bachelor’s	48.94 (50.53)	50.00 (50.51)	60.42 (49.42)	
Master’s or above	27.66 (45.22)	32.00 (47.12)	29.17 (45.93)	
N	47	50	48	145

Mean (SD) for continuous variables; mean $\times$ 100 (SD $\times$ 100) for binary.

p -values from one-way ANOVA (continuous) or chi-squared test (binary). Redistribution in USD.

4 Results

Figure 1: Effort Attribution and Merit Redistribution by Treatment Group



Notes. Left panel: mean effort-luck belief (0 = 100% luck; 10 = 100% effort), elicited after the redistribution decision. Right panel: mean USD allocated to the lower-earning worker (merit scenario). Bar heights are group means; whiskers are ± 1 SE. $N = 145$.

Figure 1 gives an overview of both main outcomes across the three treatment arms. Among the 145 spectators in the sample, 47 were told they outperformed their matched peer (above), 48 were told they underperformed (below), and 50 received no feedback. The left panel shows mean effort-luck belief and the right panel shows mean redistribution to the lower-earning worker; both outcomes display the same monotone gradient. The above group attributes the most to effort and redistributes the least; the below group is the most luck-oriented and the most redistributive; and the no-feedback group sits between the two. The most salient visual feature in both panels is the above-below gap: the distance between the two treated groups is roughly twice the distance of either group from the no-feedback anchor. By contrast, the no-feedback group sits close to the midpoint, and neither treated group differs from it by a large margin individually — a pattern consistent with both signals pushing in opposite directions from a common baseline, with the joint above-below contrast being the statistically identifiable effect.

4.1 Effort–Luck Beliefs

The left panel shows a clear monotone pattern: the above group scores highest ($\overline{EL}_{\text{above}} = 7.48$, $SD = 2.36$), followed by the no-feedback group ($\overline{EL}_{\text{nofb}} = 7.00$, $SD = 2.03$) and the below group ($\overline{EL}_{\text{below}} = 6.24$, $SD = 2.57$). Notably, all three groups score above the scale midpoint of 5, indicating that effort attribution is the modal view regardless of treatment; what the feedback shifts is the *degree* to which participants see the outcome as effort-determined. A one-way ANOVA rejects equality of group means ($p = 0.039$).

Columns 1–2 of Table 2 report OLS regressions of the effort–luck belief on treatment dummies, with the above group as the baseline. Without controls (column 1), the below group scores 1.23 points lower than the above group ($SE = 0.52$, $p < 0.05$). The no-feedback group scores 0.48 points lower ($SE = 0.46$, $p = 0.30$), roughly half the magnitude of the below-group gap and not statistically significant, consistent with the absence of any comparative signal that would activate motivated attribution. Taking the no-feedback group as a natural anchor, the point estimates suggest that positive feedback elevates effort attribution by roughly 0.48 points above this anchor, while negative feedback depresses it by approximately 0.76 points below it — both signals moving in the expected direction, though neither individual shift is large enough to be statistically significant on its own. Once controls are added (column 2), both estimates remain stable: the below-group gap widens slightly to 1.34 points ($SE = 0.53$, $p < 0.05$), and the no-feedback coefficient is -0.50 ($SE = 0.46$), confirming that the pattern does not reflect pre-existing compositional differences across groups.

These findings are consistent with the motivated-reasoning account of Kunda (1990): positive comparative signals are more readily internalised because they serve the motivated goal of maintaining a positive self-image, while negative signals invite counter-argument and discounting (Mezulis et al., 2004). Because the comparative signal is exogenously assigned and uncorrelated with actual task performance, these belief differences cannot be explained by genuine differences in effort across groups, providing experimental support for the belief-updating channel theorised by Deffains et al. (2016).

4.2 Redistribution in the Merit Scenario

The right panel shows mean redistribution following the same gradient. The above group allocates the least (\$1.91, $SD = \$1.35$), followed by the no-feedback group (\$2.12, $SD = \$1.39$) and the below group (\$2.44, $SD = \$1.38$). The one-way ANOVA yields $p = 0.179$, not significant at the omnibus level, consistent with the gap being driven mainly by the above–below contrast rather than uniform differences across all three groups.

Columns 3–4 of Table 2 report the main OLS regressions ($N = 145$). Without controls (column 3), the below group redistributes \$0.52 more than the above group ($SE = \$0.28$, $p < 0.10$), while the no-feedback group redistributes \$0.21 more ($SE = \0.28), a difference that is not statistically significant. With the full set of controls (column 4), the below-group estimate is \$0.55 ($SE = \0.31, $p < 0.10$), remaining marginally significant. The coefficient on the Colombian indicator is $-\$0.40$ ($SE = \$0.30$), negative but not statistically significant; nationality is balanced across treatment arms ($p = 0.32$), so this does not confound the treatment comparison.

The above–below gap of \$0.52–0.55 is stable across specifications and economically meaningful, representing roughly 27–29% of the above-group mean of \$1.91, and remains marginally significant in both columns.

4.3 The Belief Channel: Mediation Analysis

If self-serving attribution is the mechanism at work, then the treatment effect on redistribution should operate *through* the effort–luck belief. Table 3 tests this by adding the effort–luck belief score as a covariate to the redistribution regression.

Columns 1–2 of Table 3 use the full sample ($N = 142$, three participants missing the belief slider and one the redistribution decision). The effort–luck belief score is a strong and highly significant predictor of redistribution: each one-point increase on the 0–10 scale is associated with \$0.20 less redistributed to the lower-earning worker ($SE = \$0.05$, $p < 0.01$), both with and without additional controls. Once effort–luck belief is included, the treatment coefficients become small and statistically indistinguishable from zero: the below-group coefficient falls from \$0.52 (column 3 of Table 2) to \$0.26 ($SE = \0.28), a reduction of approximately 50 percent. The no-feedback coefficient falls from \$0.21 to \$0.09 ($SE = \$0.28$). Both results are consistent with the hypothesis that the feedback signal shifts redistribution *through* beliefs about effort and luck: once we condition on the belief, the direct effect of the signal on redistribution largely disappears.

This pattern is consistent with partial mediation by the effort–luck belief, as theorised by Deffains et al. (2016) and predicted by the accountability framework (Konow, 2000; Capellen et al., 2007): information about relative standing updates beliefs about the deservingness of inequality, and it is this updated belief — not the signal itself — that drives redistributive behaviour. Columns 3–4 of Table 3 report heterogeneity by whether the feedback was surprising given participants’ prior expectation of relative performance; insufficient variation in the surprise indicators precludes drawing conclusions from this sub-analysis with the current sample size.

Table 2: Treatment Effects on Effort Attribution and Redistribution

	(1)	(2)	(3)	(4)
Constant	7.423*** (0.348)	6.470*** (1.301)	1.915*** (0.199)	3.461*** (0.837)
No Feedback	-0.467 (0.453)	-0.476 (0.450)	0.205 (0.281)	0.172 (0.290)
Told: Below	-1.205** (0.509)	-1.301** (0.518)	0.523* (0.283)	0.549* (0.305)
Age		0.021 (0.018)		-0.008 (0.012)
Female		0.324 (0.405)		-0.256 (0.237)
Colombian		0.413 (0.533)		-0.402 (0.296)
Task score		-0.007 (0.009)		-0.005 (0.006)
Bachelor's degree		0.369 (0.549)		-0.363 (0.365)
Master's or above		0.461 (0.597)		-0.372 (0.404)
N	145	145	145	145
R^2	0.044	0.110	0.024	0.077
Controls (age, female, score, Colombian, edu.)	No	Yes	No	Yes

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$

Notes. OLS with HC3-robust SEs in parentheses. Baseline: Told Above. Cols. (1)–(2): effort–luck belief (0 = luck, 10 = effort). Cols. (3)–(4): USD redistributed to lower-earning worker (merit scenario). Controls: age, female indicator, Colombian indicator, task score, and education (Less than bachelor = baseline). * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 3: Mediation Analysis

	(1)	(2)	(3)	(4)
Constant	3.433*** (0.432)	4.717*** (0.796)	2.179*** (0.143)	2.950** (1.203)
No Feedback	0.110 (0.274)	0.080 (0.280)		
Told: Below	0.276 (0.274)	0.297 (0.292)		
Effort–luck belief	−0.204*** (0.051)	−0.194*** (0.053)		
Age		−0.004 (0.011)		−0.001 (0.016)
Female		−0.193 (0.229)		0.098 (0.343)
Colombian		−0.321 (0.291)		−0.416 (0.381)
Task score		−0.007 (0.005)		−0.002 (0.009)
Bachelor’s degree		−0.292 (0.339)		−0.164 (0.447)
Master’s or above		−0.283 (0.378)		−0.503 (0.552)
<i>N</i>	145	145	95	95
<i>R</i> ²	0.140	0.174	0.000	0.040
Controls (age, female, score, Colombian, edu.)	No	Yes	No	Yes

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$

Notes. OLS with HC3-robust SEs in parentheses. Dependent variable: USD redistributed (merit scenario). Cols. (1)–(2): full sample; includes effort–luck belief to assess mediation; compare treatment coefficients with Table 2 cols. (3)–(4). Cols. (3)–(4): relative-feedback group only. “Positively surprised”: told above but predicted below. “Negatively surprised”: told below but predicted above. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

5 Conclusion

This paper provides experimental evidence that relative performance feedback shapes beliefs about the role of effort in generating inequality and, through that channel, redistribution preferences. Participants told they outperformed their peer attribute significantly more of the task outcome to effort than those told they underperformed (gap of 1.23 points on a 0–10 scale, $p < 0.05$), and redistribute less in the merit scenario (gap of \$0.52–0.55 relative to the below group, $p < 0.10$). No significant redistribution effect is found relative to the no-feedback group, though the above–below gap is statistically significant and consistent with a self-serving attribution channel.

Taken together, these findings are consistent with a self-serving attribution account of redistribution preferences: learning one outperformed a peer reinforces the belief that inequality reflects effort, which in turn reduces willingness to redistribute in settings where merit is salient. That the effect is most clearly identified by contrasting participants who received positive versus negative comparative signals — rather than feedback versus no feedback — suggests the *direction* of the relative signal, not mere exposure to comparison, is the key mechanism. This asymmetry mirrors the well-documented pattern in the psychology of self-serving attribution (Mezulis et al., 2004) and has implications for the design of performance information systems.

These findings have implications for the study of inequality and redistribution in a world where performance comparisons are increasingly salient — on professional platforms, social networks, and in educational settings. The self-serving updating documented here suggests that rising relative visibility of performance rankings could systematically reduce support for redistribution among those who “win” such comparisons, potentially contributing to the polarisation of redistribution preferences.

Future work should (i) increase sample size to improve power, particularly for the heterogeneity and mediation analyses; (ii) test whether the effect persists when participants receive absolute score information alongside relative feedback, to disentangle comparative and informational channels; and (iii) conduct a formal mediation analysis to quantify the share of the redistribution effect that operates through the effort-attribution belief measure.

References

- Alesina, Alberto, and Paola Giuliano. 2011. “Preferences for Redistribution.” In *Handbook of Social Economics*, Vol. 1A, edited by Jess Benhabib, Alberto Bisin, and Matthew O. Jackson, 93–131. Amsterdam: North-Holland.
- Alesina, Alberto, and George-Marios Angeletos. 2005. “Fairness and Redistribution.” *American Economic Review*, 95(4): 960–980.
- Almås, Ingvild, Alexander W. Cappelen, and Bertil Tungodden. 2020. “Cutthroat Capitalism versus Cuddly Socialism: Are Americans More Meritocratic and Efficiency-Seeking than Scandinavians?” *Journal of Political Economy*, 128(5): 1753–1788.
- Babcock, Linda, and George Loewenstein. 1995. “Explaining Bargaining Impasse: The Role of Self-Serving Biases.” *Journal of Economic Perspectives*, 9(1): 109–126.
- Bénabou, Roland, and Jean Tirole. 2006. “Belief in a Just World and Redistributive Politics.” *Quarterly Journal of Economics*, 121(2): 699–746.
- Konow, James. 2000. “Fair Shares: Accountability and Cognitive Dissonance in Allocation Decisions.” *American Economic Review*, 90(4): 1072–1091.
- Cappelen, Alexander W., Astri Drange Hole, Erik Ø. Sørensen, and Bertil Tungodden. 2007. “The Pluralism of Fairness Ideals: An Experimental Approach.” *American Economic Review*, 97(3): 818–827.
- Cappelen, Alexander W., James Konow, Erik Ø. Sørensen, and Bertil Tungodden. 2013. “Just Luck: An Experimental Study of Risk-Taking and Fairness.” *American Economic Review*, 103(4): 1398–1413.
- Cassar, Lea, and Arnd H. Klein. 2019. “A Matter of Perspective: How Failure Shapes Distributive Preferences.” *Management Science*, 65(11): 5050–5064.
- Deffains, Bruno, Romain Espinosa, and Christian Thöni. 2016. “Political Self-Serving Bias and Redistribution.” *Journal of Public Economics*, 134: 67–74.
- Hastorf, Albert H., and Hadley Cantril. 1954. “They Saw a Game: A Case Study.” *Journal of Abnormal and Social Psychology*, 49(1): 129–134.
- Mezulis, Amy H., Lyn Y. Abramson, Janet S. Hyde, and Benjamin L. Hankin. 2004. “Is There a Universal Positivity Bias in Attributions? A Meta-Analytic Review of Individual,

- Developmental, and Cultural Differences in the Self-Serving Attributional Bias.” *Psychological Bulletin*, 130(5): 711–747.
- Kunda, Ziva. 1990. “The Case for Motivated Reasoning.” *Psychological Bulletin*, 108(3): 480–498.
- Miller, Dale T., and Michael Ross. 1975. “Self-Serving Biases in the Attribution of Causality: Fact or Fiction?” *Psychological Bulletin*, 82(2): 213–225.
- Piketty, Thomas. 2014. *Capital in the Twenty-First Century*. Cambridge, MA: Harvard University Press.
- Piketty, Thomas. 1995. “Social Mobility and Redistributive Politics.” *Quarterly Journal of Economics*, 110(3): 551–584.